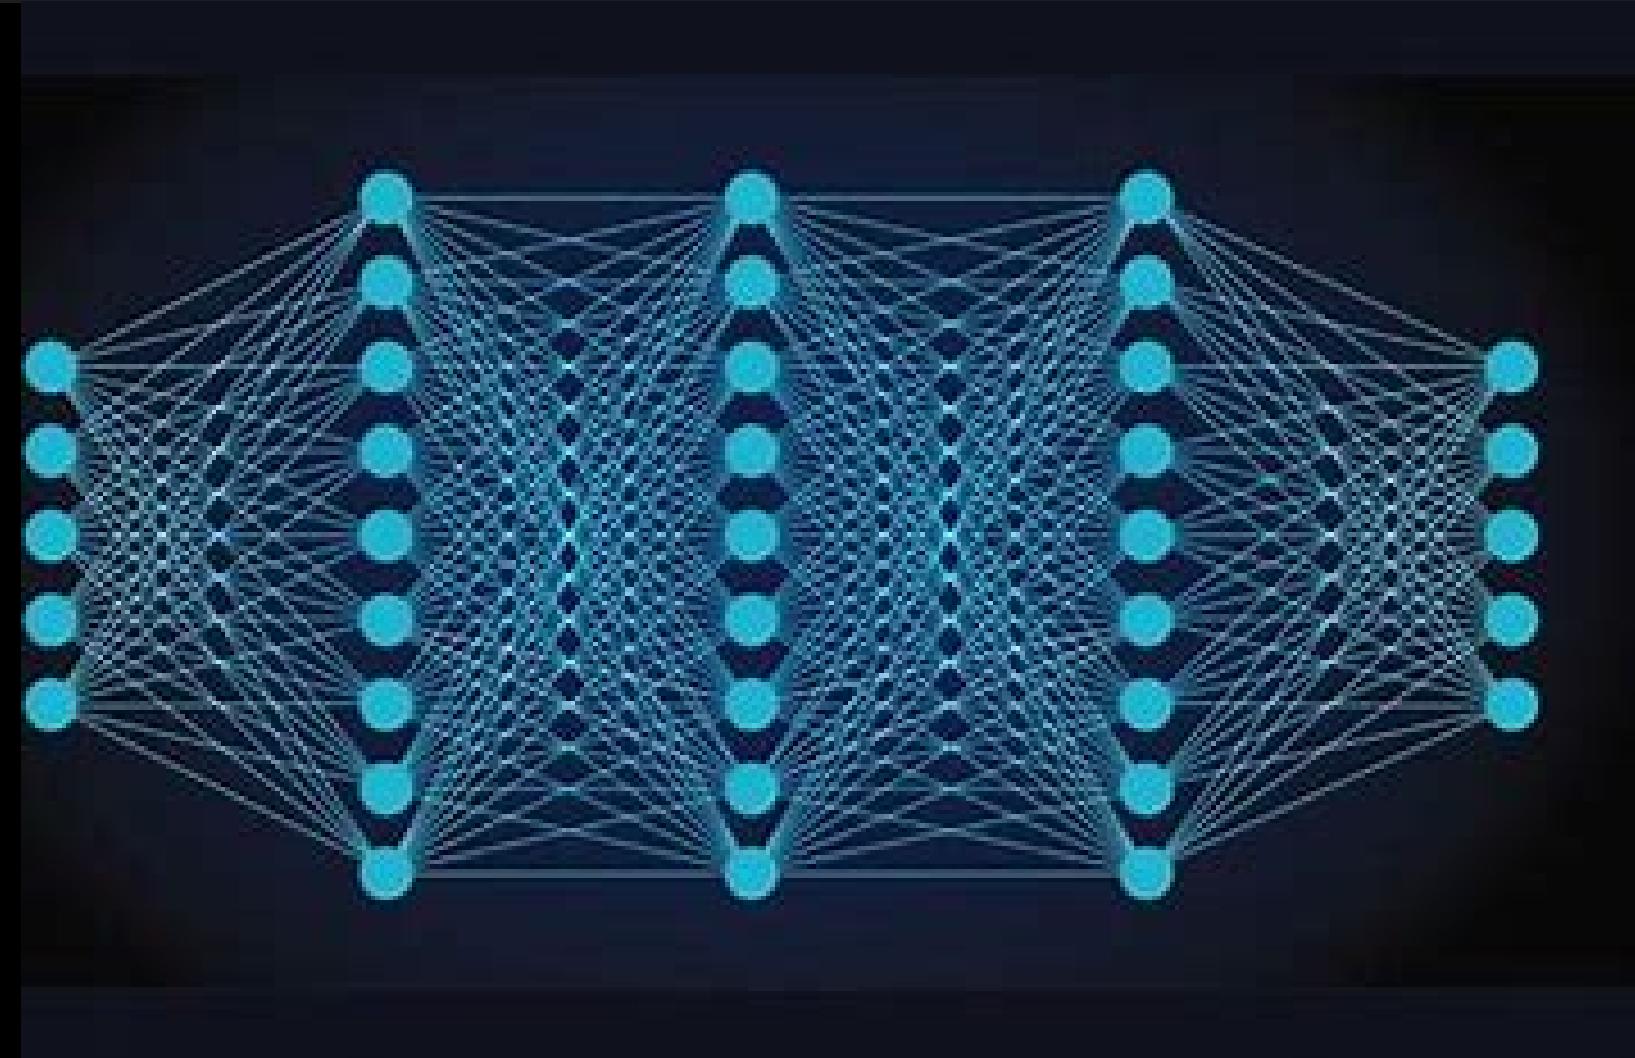




Artificial Intelligence 101

Mid - Term Evaluation

- Harnessing the power of interconnected layers to unravel mysteries in AI, from perceiving the unseen to understanding the unsaid.



Watt Wizards
By Group - 6

Intro To ML

How does it work?



We give the data to our model



And the model makes predictions



TYPES

Supervised Learning

The model learns from labeled data to predict outcomes.

Examples:

Regression: Predict house prices.
Classification: Email spam detection.

Semi-Supervised Learning

The model learns from a mix of labeled and unlabeled data

Example:

Text categorization with partially labeled documents.

Unsupervised Learning

The model finds hidden patterns or groupings in data without labels.

Examples:

Clustering: Customer segmentation.
Dimensionality Reduction: PCA for visualization.

Maximum Likelihood Estimation (MLE)

A statistical method to estimate the parameters of a model by maximizing the likelihood function.

Goal: Find parameter values (θ) that maximize the likelihood of observed data (X)

Steps-

Define a probabilistic model:

Example: Assume data follows a Normal distribution.

Maximize the likelihood (or log-likelihood)

Solve for μ and σ :

- $\mu = \frac{1}{n} \sum_{i=1}^n x_i$ (mean).
- $\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$ (variance).

Write the likelihood function

Ex:- For Normal distribution:

$$L(\mu, \sigma | X) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$$

CLASSIFICATION IN ML

One of the fundamental task in ML which involves predicting the category or the class label of an input based on its features

REQUIREMENTS?

Labelled Data

Feature Selection and Representation

Algo(like SVM ,logistic regression) and Evaluation matrices(like Euclidean matrice)

Linear Vs Non-Linear Classifiers

Linear classifiers assumes that classes can be separated by a straight line or a hyperplane while non-linear creates a more complex curve and do not assume linear boundaries.

K-Nearest Neighbour

Non-Linear Classifier which classifies the input data by looking at k-nearest data points to it from the training data. The algorithm uses distance matrices (like Euclidean Distance) to measure the distance between data points.

Support Vector Machines(SVMs)

Linear Classifier which aims to find the optimum margin line(or Hyperplane) that separates data points of different class. Unlike Knn it does not store the training data.

How SVM Works?

Finds the hyperplane $w \cdot x + b = 0$, that separates the data, maximizing the margin

TYPES

Hard SVM:- Assumes data is perfectly linearly separable and finds a hyperplane with no misclassification.

Soft SVM:- Allows for misclassification and finds a hyperplane with maximum margin but with a trade off between maximizing the margin minimizing misclassification.

For separable data, the optimisation and the constraints are :-

$$\begin{aligned} & \min_{w,b} \frac{1}{2} \|w\|^2 \\ \text{s.t. } & y^{(i)}(w^T x^{(i)} + b) \geq 1, \quad i = 1, \dots, n \end{aligned}$$

Intro to Random variable

In probability, a real-valued function, defined over the sample space of a random experiment, is called a **random variable**.

Example:

Consider the random experiment of tossing a coin $S=\{H,T\}$
Define Random variable $X(H)=1, X(T)=0$

Random Variables

Continuous
Random
Variables

Discrete
Random
Variables



Discrete
variable

Countable support

Probability mass function

Probabilities assigned to single values

Each possible value has strictly positive probability

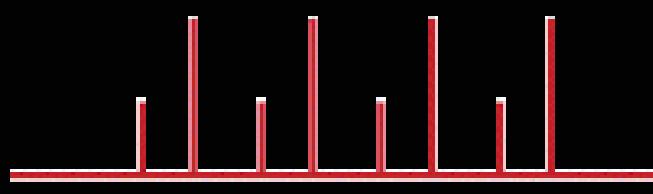
Continuous
variable

Uncountable support

Probability density function

Probabilities assigned to intervals of values

Each possible value has zero probability



PDF Transformation - CDF Approach

Given!!!

C.R.V. X with pdf $f_X(x)$
and $Y = u(X)$, $u(X)$ need
to be invertible



To find
PDF of Y



CDF of Y

$$f_Y(y) = \frac{d}{dy} F_Y(y)$$

PDF of Y

$$f_Y(y) = \frac{d}{dy} F_X(u^{-1}(y))$$

$$f_Y(y) = f_X(u^{-1}(y)) \cdot \frac{d}{dy} u^{-1}(y)$$

$$f_Y(y) = f_X(u^{-1}(y)) \cdot \left| \frac{du^{-1}(y)}{dy} \right|$$

$$F_Y(y) = P[Y \leq y]$$

$$F_Y(y) = P[u(X) \leq y]$$

$$F_Y(y) = P[X \leq u^{-1}(y)]$$

$$F_Y(y) = F_X(u^{-1}(y))$$

$$f_Y(y) = f_X(u^{-1}(y)) \cdot \left| \frac{du^{-1}(y)}{dy} \right|$$

Result

Central Limit Theorem

1 Lindeberg-Levy Theorem

Theorem 1 Let $\{X_n\}$ be a sequence of i.i.d. r.v.s such that $\mathbb{E}[X_k] = 0$, $\text{Var}[X_k] = \sigma^2 < \infty$ and define $S_n = \sum_{k=1}^n X_k$. Then as $n \rightarrow \infty$,

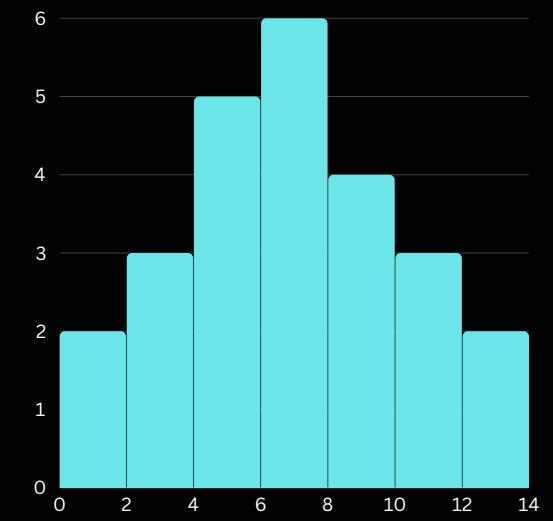
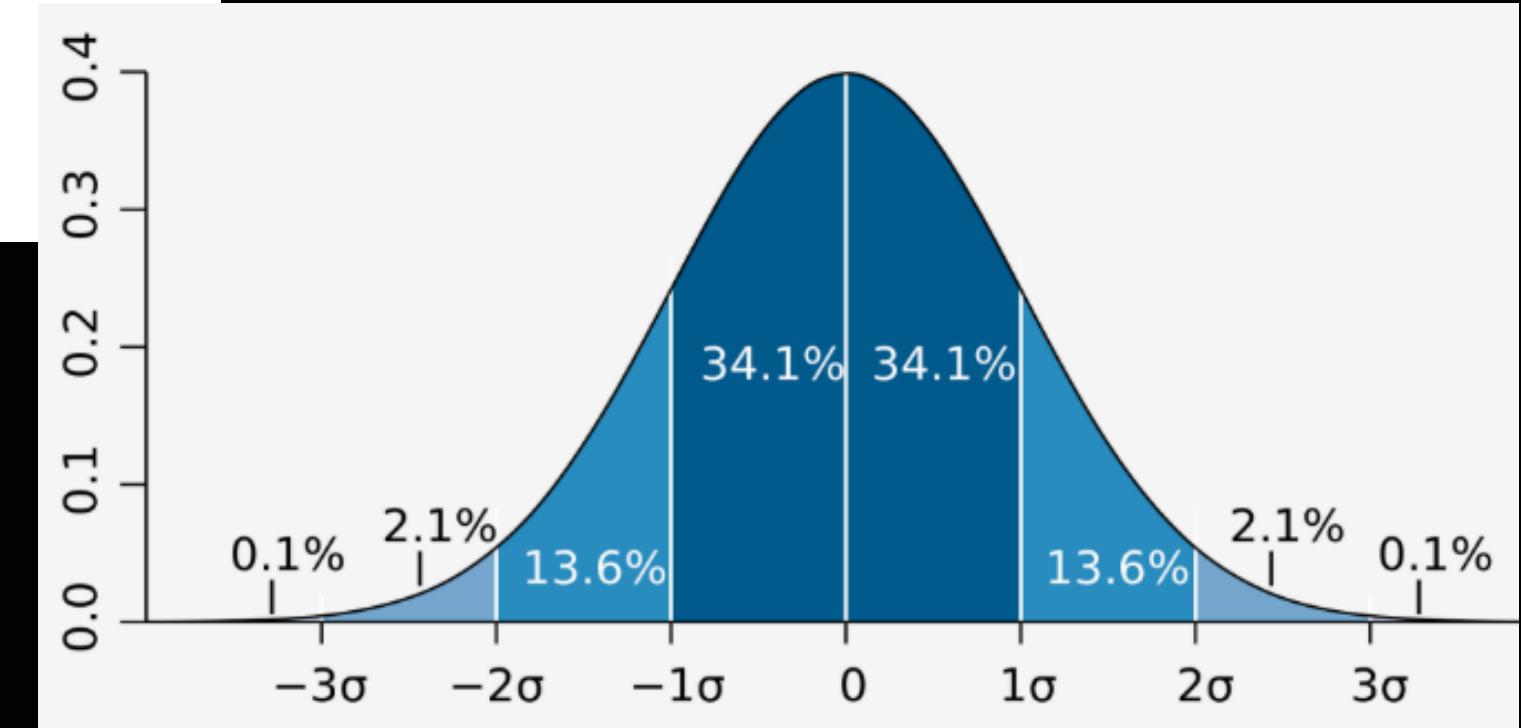
$$\frac{S_n}{\sigma\sqrt{n}} \xrightarrow{d} \tau \text{ where } \tau \sim N(0,1).$$

All X_i 's
are independent
of each other.



2 Each X_i is drawn from
the same distribution.

3 Finite Variance
($0 < \text{Var}(X_i) < \infty$).



Cumulative Distribution Function:

Function that describes the probability that a random variable is less than or equal to a given value. It is usually represented as $F(x)=P[X \leq x]$.

Properties:

Never decreasing
 $F(x) \leq F(y)$ for $x < y$

F is right continuous
 $F(x) = F(x^+)$

$$\lim_{x \rightarrow -\infty} F(x) = 0$$

$$\lim_{x \rightarrow \infty} F(x) = 1$$

$f(x)$ must satisfy these 2 conditions. For discrete integration summation take place

. We require:

$$f(x) \geq 0 \text{ for all } x.$$

. We require:

$$\int_{-\infty}^{\infty} f(x) dx = 1.$$

Relationship between pdf and cdf

The cdf is obtained from the pdf through integration:

$$F(x) = P(X \leq x) = \int_{-\infty}^x f(t) dt \quad \text{for all } x.$$

The pdf is obtained from the cdf through differentiation:

$$f(x) = F'(x).$$

Expectation (mean):

E(X)pectation of a Discrete Random Variable

$$E(X) = \sum_x xP(x)$$

E(X)pectation of a Continuous Random Variable

$$E(X) = \int_{-\infty}^{\infty} xP(x)dx$$



Properties of Expectations

1. $E[c] = c$
2. $E[aX + b] = aE[X] + b$
3. $\sigma^2(ax + b) = a^2\sigma^2$
4. $E[g(x)] = \int g(x)dF(x)$

$$\sigma^2 = \sum (Xi - u)^2/n$$

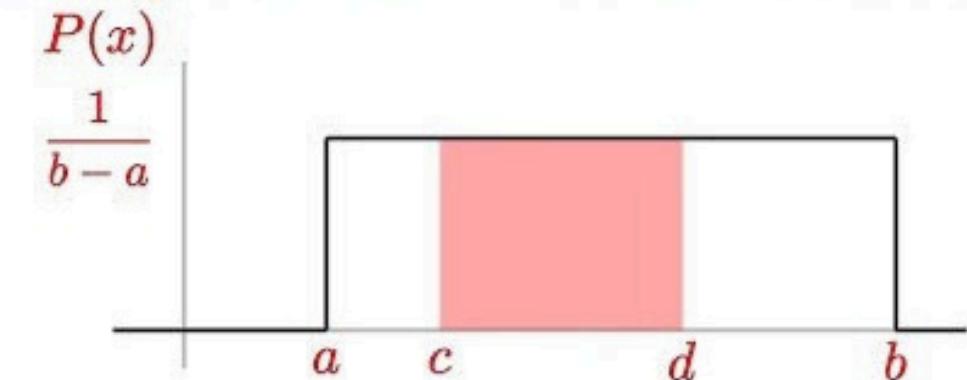
- n is the total number of observation

$$\sigma = \sqrt{E(X^2) - (E(X))^2}$$

- σ denotes variance

- $u = E(x)$

Uniform Distribution

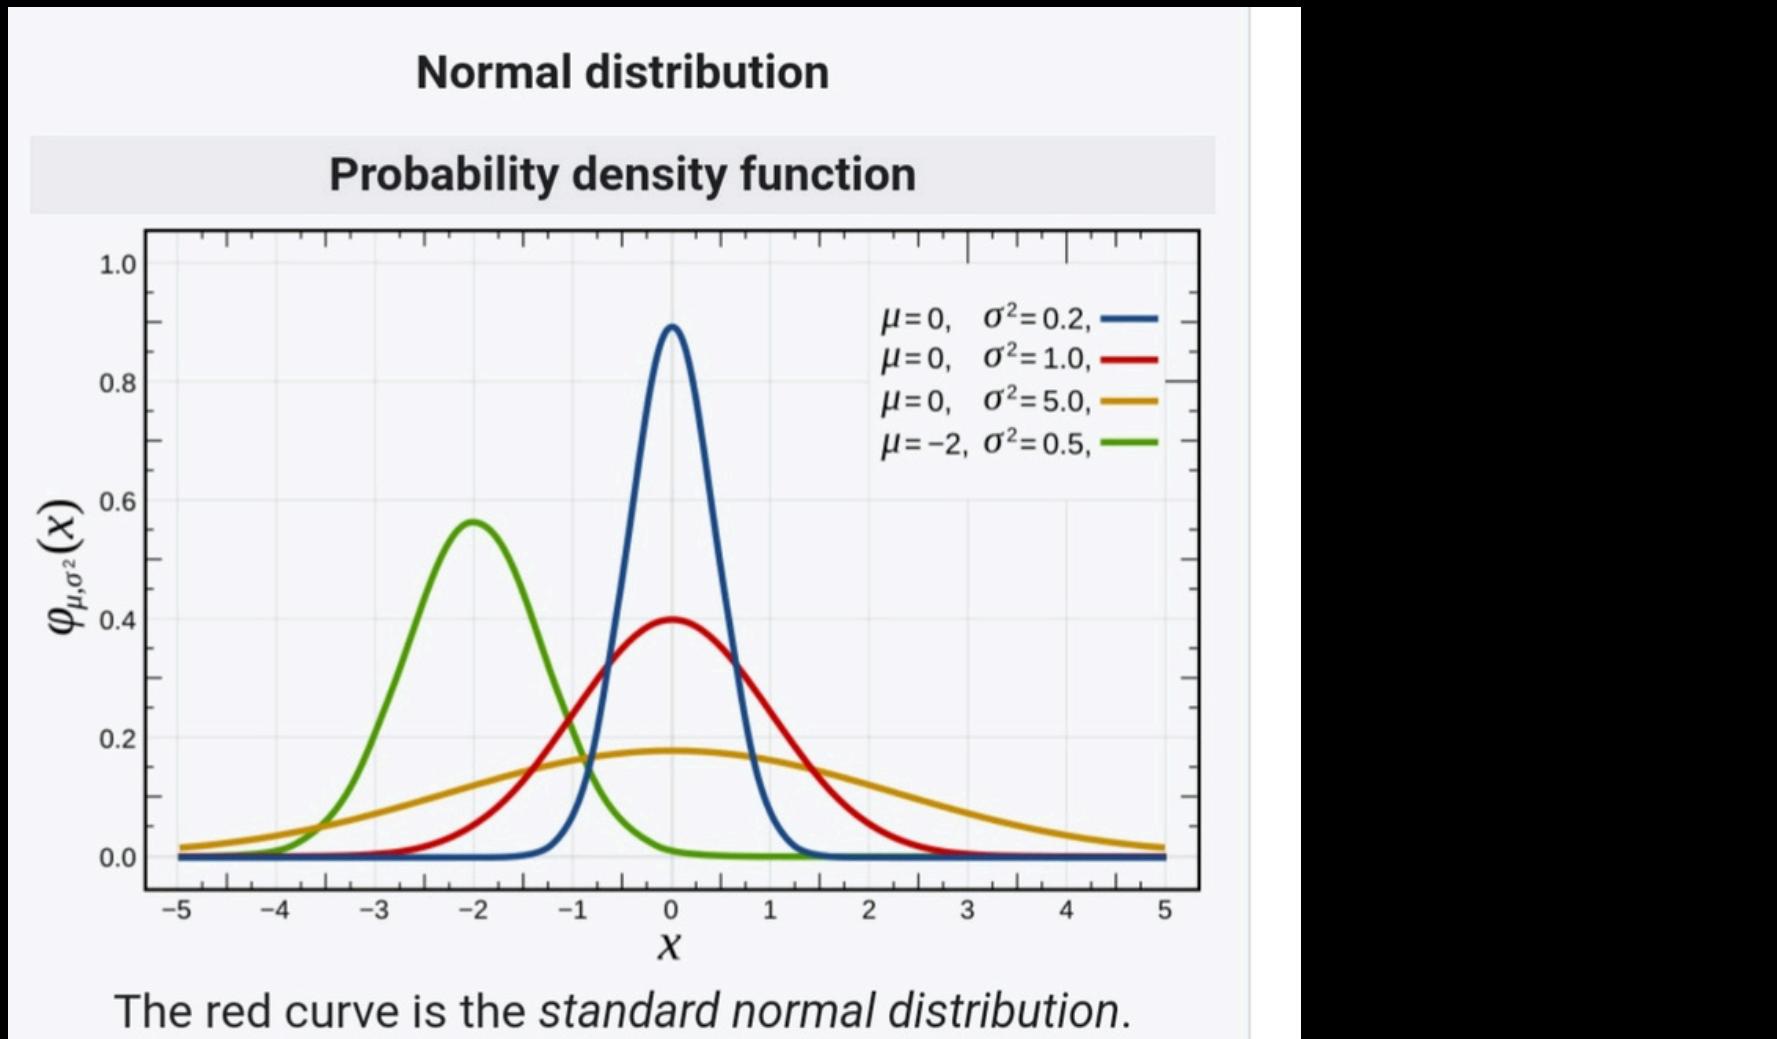


$$\text{Mean: } \mu = \frac{a+b}{2}$$

$$\text{S.D.: } \sigma = \sqrt{\frac{(b-a)^2}{12}} \quad P(c \leq X \leq d) = \frac{d-c}{b-a}$$

Probability Density Formula for Normal Distribution

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x - \mu)^2}{2\sigma^2}}$$



- In standard normal distribution $\mu=0$ and $\sigma=1$
Symmetric about origin

The CDF of the standard normal distribution is denoted by the Φ function:

$$\Phi(x) = P(Z \leq x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x \exp \left\{ -\frac{u^2}{2} \right\} du.$$

Marginal Distribution Function

Continuous Random Variable

If the joint distribution of the random variables X , Y is $F[x, y]$, then

the marginal distribution of X is

$$F_X(x) = \int_{-\infty}^x \int_{-\infty}^{\infty} f(x, y) dy dx$$

the marginal distribution of Y is

$$F_Y(y) = \int_{-\infty}^y \int_{-\infty}^{\infty} f(x, y) dx dy$$

Convolutional Neural Networks

Convolution

$$\begin{array}{|c|c|c|c|} \hline 3 & 5 & 4 & 2 \\ \hline 6 & 2 & 2 & 4 \\ \hline 7 & 3 & 5 & 1 \\ \hline 3 & 2 & 3 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -1 & -1 \\ \hline \end{array} = \begin{array}{|c|c|} \hline -3 & -2 \\ \hline 2 & 2 \\ \hline \end{array}$$

The filter slides over the input and computes the element-wise product between the filter and the portion of the input it covers.

Extracts local features and reduces dimensions.

Padding

Padding involves adding extra pixels (usually zeros) around the input image to control the spatial dimensions of the output after convolution.

Stride

Stride refers to the step size the filter moves while scanning the input data. Larger stride values result in filter moving more quickly across the input.

Pooling

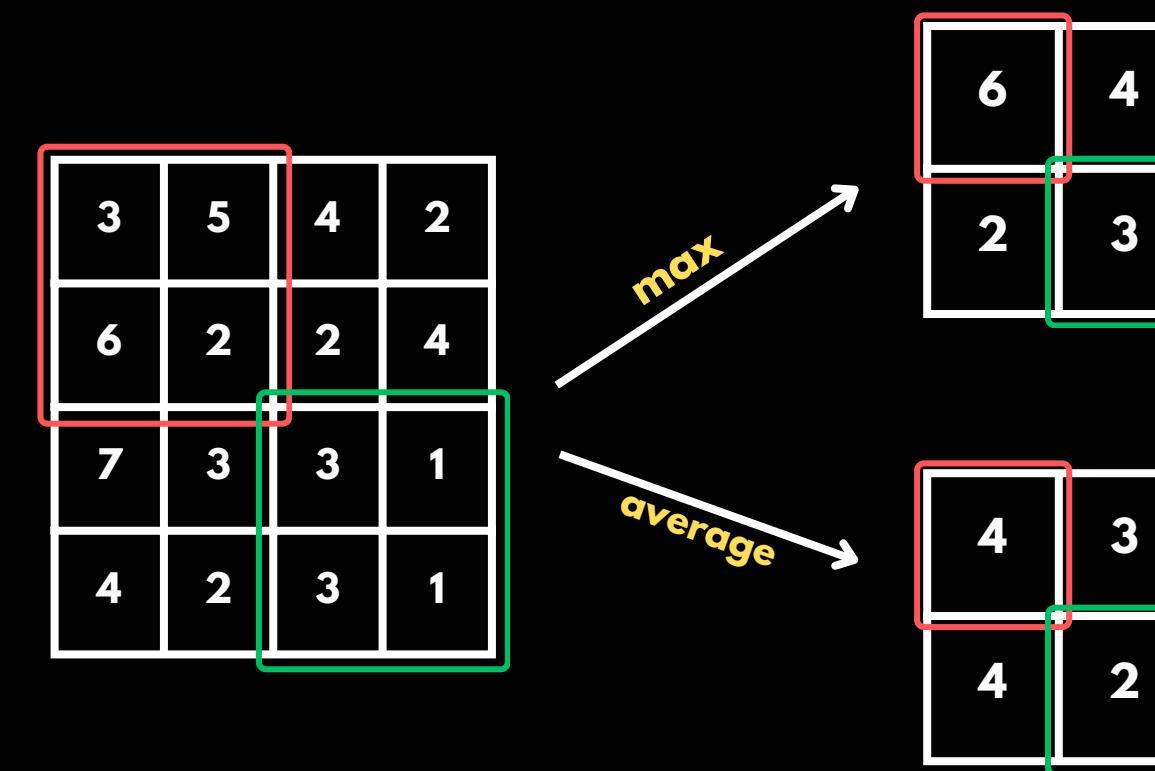
Pooling is used to reduce the size of an image or feature map while keeping important information. This helps the network learn more efficiently by focusing on the most important features and reducing unnecessary details.

Max Pooling

For each small region (like a 2x2 square), max pooling takes the largest number.

Average Pooling

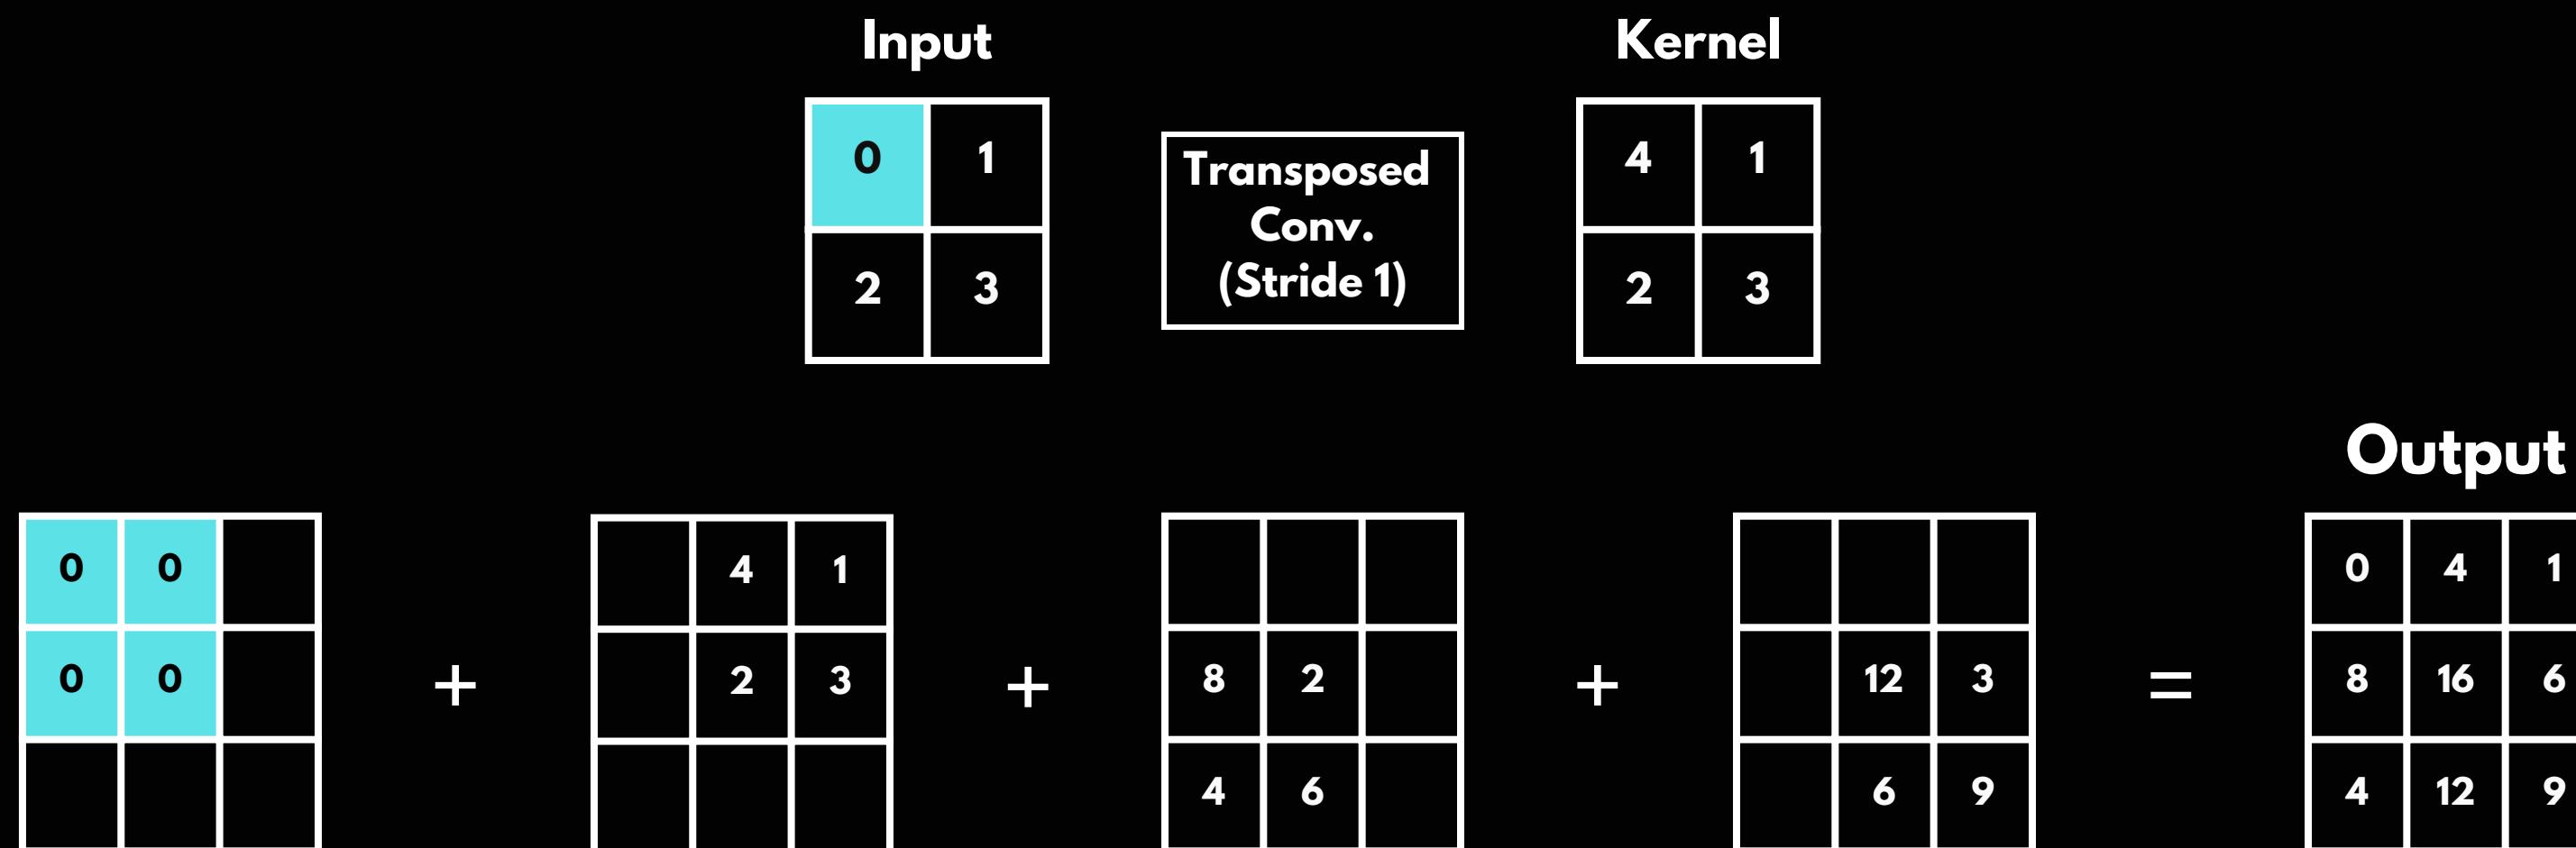
For each small region (like a 2x2 square), average pooling takes the average.



Transposed Convolution

Transposed convolution can be thought of as the reverse of regular convolution. While convolution is designed to reduce spatial dimensions, transposed convolution is used to expand these dimensions. It "reverses" the effect of downsampling by generating more detailed information in a larger spatial area, allowing the network to recreate higher-resolution features from lower-dimensional representations.

It is often used for upsampling smaller feature maps, such as in image-to-image translation tasks, where the goal is to take a compressed or lower-dimensional representation and generate a high-resolution output.

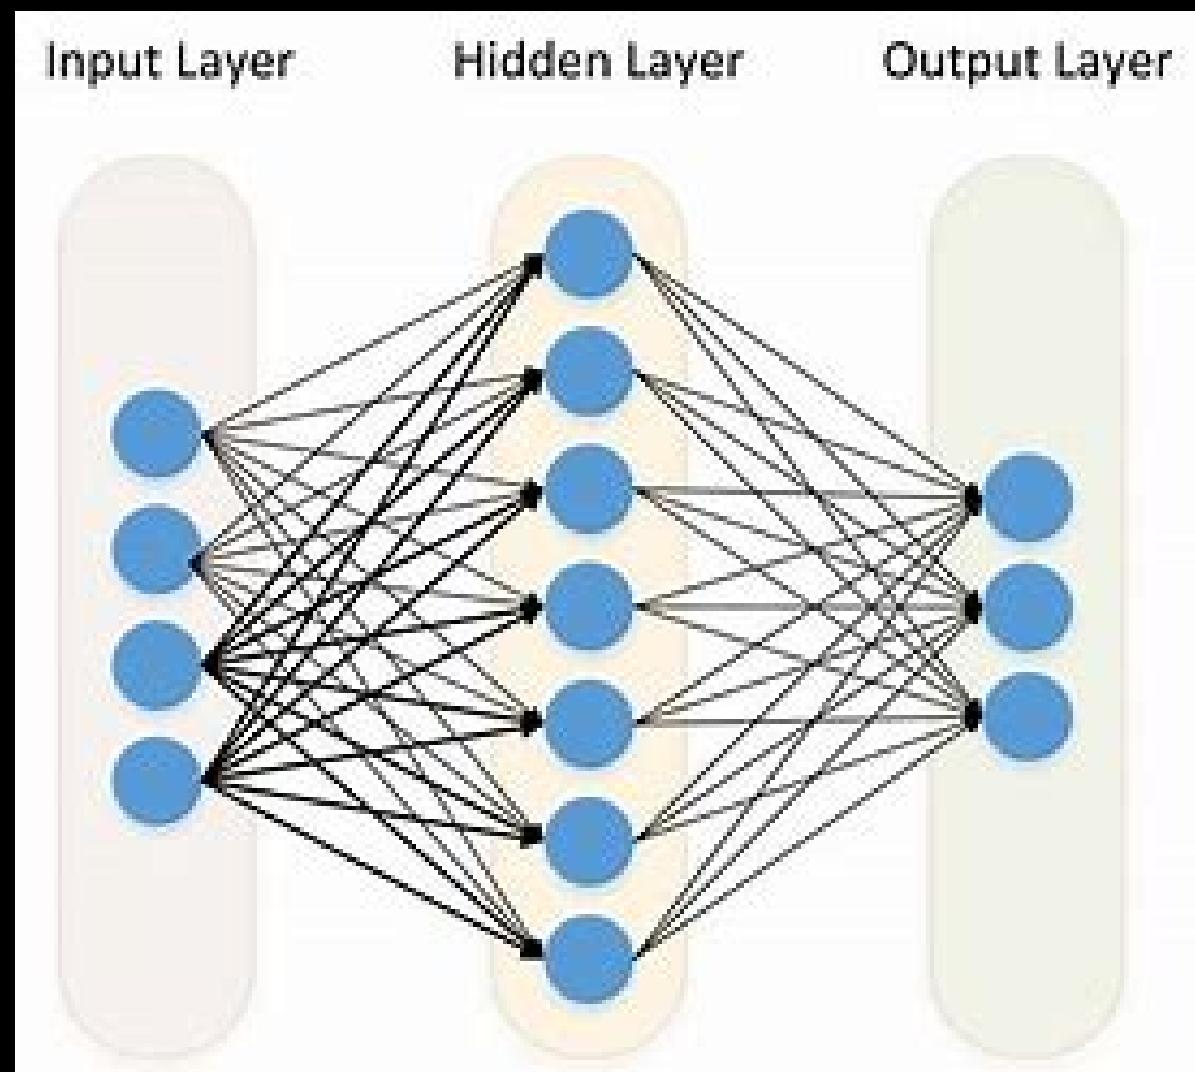


Neural Network

NN Models are computational models inspired by the brain's interconnected neurons for information processing.

- They consist of three main layers: input, hidden and output layers.
- NN are widely employed in image recognition, speech processing and various ML tasks.

Training involves adjusting weights and biases among the neurons and using feature maps to easily analyze the data and the use of an activation function to achieve our result.



Backpropagation

- Purpose: Backpropagation is used to train neural networks by optimizing weights and biases to minimize the error.

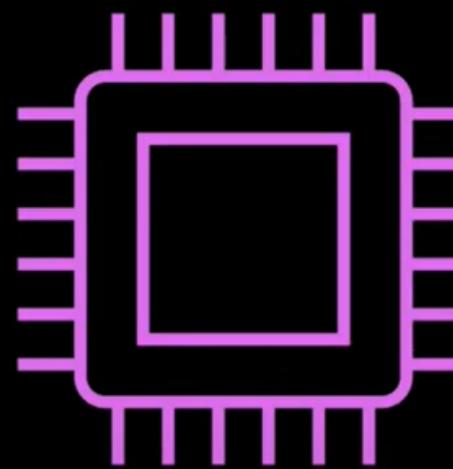
Process overview

Forward Pass
Input flows through the network.

Loss Calculation
Compare predictions with actual values.

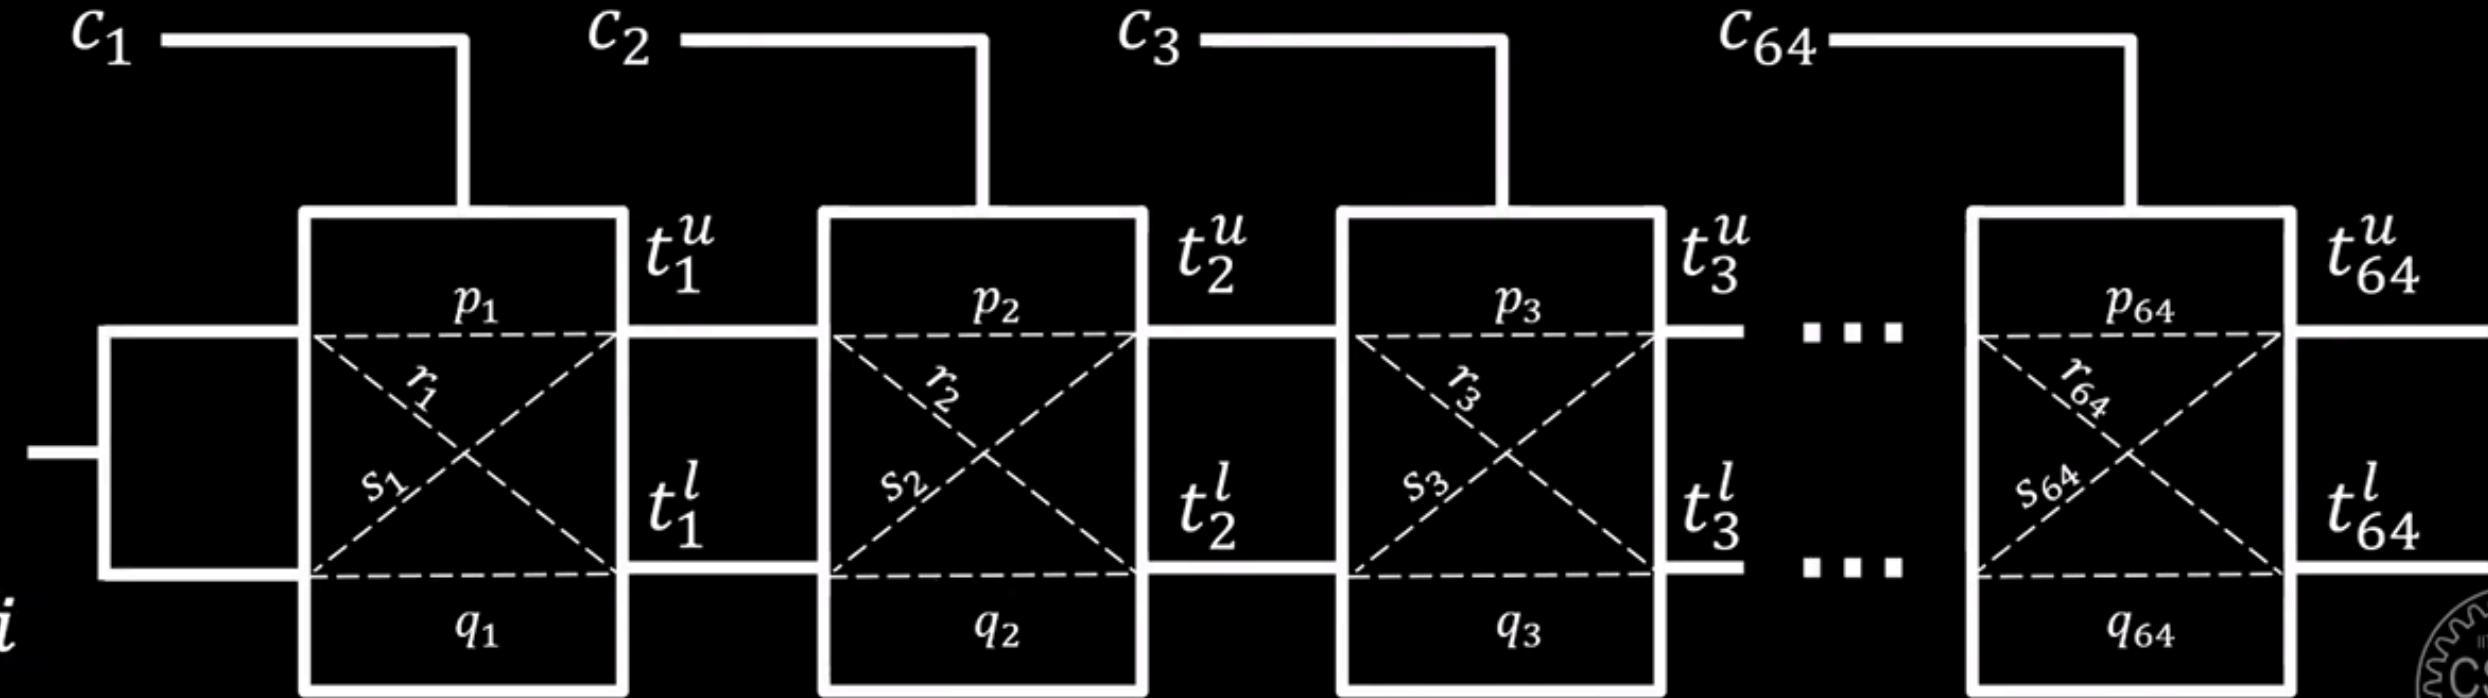
Backward Pass
Calculate gradients using the chain rule.

Weight Update
Adjust weights and biases using gradients and learning rate.



PUFs - Physically Unclonable Functions.

$$\Delta_i = t_i^u - t_i^l$$
$$\Delta_i = d_i \cdot \Delta_{i-1} + \alpha_i \cdot d_i + \beta_i$$



We can keep going on recursively

$$\Delta_1 = \alpha_1 \cdot d_1 + \beta_1 \text{ (since } \Delta_0 = 0)$$

$$\Delta_2 = \Delta_1 \cdot d_2 + \alpha_2 \cdot d_2 + \beta_2 : \text{now plugin value of } \Delta_1 \text{ to get}$$

$$\Delta_2 = \alpha_1 \cdot d_2 \cdot d_1 + (\alpha_2 + \beta_1) \cdot d_2 + \beta_2$$

$$\Delta_3 = \alpha_1 \cdot d_3 \cdot d_2 \cdot d_1 + (\alpha_2 + \beta_1) \cdot d_3 \cdot d_2 + (\alpha_3 + \beta_2) \cdot d_3 + \beta_3$$

$$\begin{aligned}\Delta_4 = & \alpha_1 \cdot (d_4 \cdot d_3 \cdot d_2 \cdot d_1) + (\alpha_2 + \beta_1) \cdot (d_4 \cdot d_3 \cdot d_2) \\ & + (\alpha_3 + \beta_2) \cdot (d_4 \cdot d_3) + (\alpha_4 + \beta_3) \cdot (d_4) + \beta_4\end{aligned}$$



Do you see a pattern emerging here too!

It turns out that

$$\Delta_{64} = \mathbf{w}^T \mathbf{x} + b$$

kuch yadd
aaya !!!!



Yess!!!, ye
to Linear
model hai



where $\mathbf{x} = [x_1, \dots, x_{64}]$, $\mathbf{w} = [w_1, \dots, w_{64}] \in \mathbb{R}^{64}$

$$\begin{aligned} x_i &= d_i \cdot d_{i+1} \cdot \dots \cdot d_{64} \\ &= (1 - 2c_i) \cdot (1 - 2c_{i+1}) \cdot \dots \cdot (1 - 2c_{64}) \end{aligned}$$

$$w_1 = \alpha_1$$

$$w_i = \alpha_i + \beta_{i-1} \text{ (for } i = 2, 3, \dots, 64\text{)}$$

$$b = \beta_{64}$$

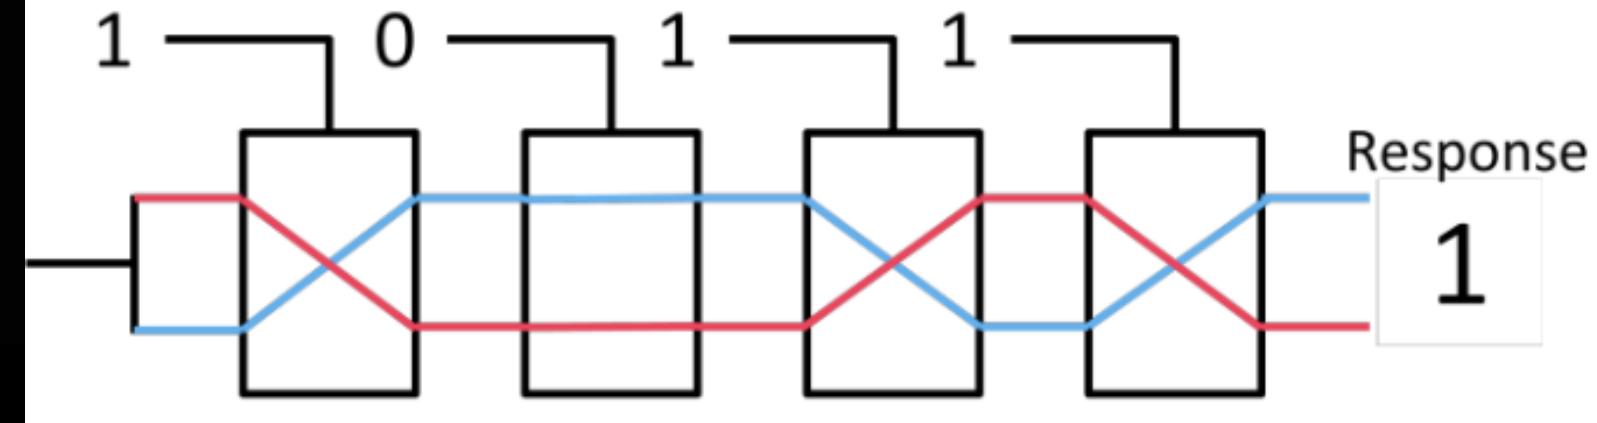
Thus, response is simply $\frac{\text{sign}(\mathbf{w}^T \mathbf{x} + b) + 1}{2}$

if $\Delta_{64} < 0 \rightarrow \text{0} \text{, else } \text{1}$



A Simple Arbiter PUF.

Challenge: 1011



$$d_i \stackrel{\text{def}}{=} (1 - 2c_i)$$

$$\alpha_i \stackrel{\text{def}}{=} (p_i - q_i + r_i - s_i)/2$$

$$\beta_i \stackrel{\text{def}}{=} (p_i - q_i - r_i + s_i)/2$$

Thank You!

Mentees:-

Kaushal Mehra, 230546
Rohit Singh, 230870
Sarthak Shekhar, 230929
Nitesh Kumar, 230707
Karan Kumar, 230533
Naman Patidar, 230679
Krishna Agarwal, 230574
Shital Niras, 230967
Vaibhav Chauhan, 231113
Akshat Dhote, 230095
Dishank Kumawat, 230377

Mentors:-

K.N. Joshua
Nikhil Jain
Aravind